



Assignment No. 6

Title:

Data Security and Privacy Assessment.

Aim:

To identify potential data security risks in big data environments and implement privacy-preserving measures to protect sensitive information. The objective is to ensure secure storage, processing, and sharing of large datasets while maintaining confidentiality, integrity, and compliance with data protection regulations

Objective:

The main objective of this lab is to understand the principles of data security and privacy in big data environments. This experiment helps students identify potential security risks, including unauthorized access, data breaches, and leakage of sensitive information. Furthermore, students will implement privacy-preserving techniques such as data anonymization, encryption, and access controls to protect sensitive datasets. Through this process, they will gain practical knowledge of securing big data, ensuring compliance with data protection regulations, and maintaining confidentiality, integrity, and availability of information in large-scale data systems.

Software Requirements:

1. Operating System: Windows / Linux / macOS
2. Python 3.x
3. Apache Spark (PySpark)
4. Apache Hadoop (optional, for HDFS storage)
5. Jupyter Notebook / VS Code / PyCharm
6. Required Python libraries: pandas, numpy, scikit-learn, matplotlib, cryptography.
7. Sample large dataset containing sensitive or private information (CSV file or stored in HDFS)

Outcome:

After completing this lab, students will be able to:

1. Identify and assess potential data security risks in big data environments.
2. Implement privacy-preserving techniques such as data anonymization, masking, and encryption.
3. Apply access control and authorization measures to protect sensitive datasets.
4. Understand best practices for ensuring data confidentiality, integrity, and compliance with privacy regulations.



PCET's
Pimpri
Chinchwad
University
Learn | Grow | Achieve

PIMPRI CHINCHWAD UNIVERSITY School of Engineering and Technology



5. Gain hands-on experience in securing and managing sensitive big data in distributed systems.

Theory:

In big data applications, vast amounts of data are collected from multiple sources such as social media platforms, business transactions, IoT devices, healthcare systems, and online services. These datasets often contain sensitive and personal information, including financial records, personal identities, medical details, location data, and confidential organizational data. Because big data environments involve distributed storage systems and large-scale processing, they introduce significant security and privacy challenges. Unauthorized access, data breaches, cyberattacks, insider threats, and leakage of confidential information are some of the major risks associated with big data systems.

Traditional security mechanisms were mainly designed for centralized databases and smaller datasets. However, big data frameworks such as Apache Hadoop and Apache Spark process data across clusters of multiple machines, which increases the attack surface and requires more advanced security mechanisms. In distributed environments, data is continuously transmitted between nodes, stored in multiple locations, and processed by various applications. This makes it essential to implement strong security controls and privacy-preserving techniques to protect sensitive information.

Data security in big data focuses on protecting information from unauthorized access, modification, or loss, while data privacy ensures that personal information about individuals is not disclosed or misused. Organizations handling large-scale datasets must implement security strategies at every stage of the data lifecycle, including data collection, storage, processing, sharing, and deletion. These strategies help maintain the three key principles of information security: confidentiality, integrity, and availability (CIA triad). Ensuring data security and privacy in big data involves multiple strategies:

1. Data Encryption – Protects sensitive data at rest and in transit using cryptographic techniques so that unauthorized users cannot read it. Encryption converts plain text data into unreadable ciphertext using encryption algorithms and secret keys. Only authorized users with the correct decryption keys can access the original data. Encryption is commonly applied when storing data in distributed storage systems and when transmitting data



between servers in a cluster.

2. **Access Control and Authentication** – Restricts who can access or modify the data using roles, permissions, and secure authentication mechanisms. Access control ensures that only authorized users or applications can interact with specific datasets. Role-based access control (RBAC) allows organizations to assign different permissions to users based on their roles. Authentication methods such as passwords, biometric verification, tokens, and multi-factor authentication provide additional layers of protection.
3. **Data Anonymization and Masking** – Removes or masks personally identifiable information (PII) such as names, phone numbers, email addresses, and identification numbers. Techniques such as pseudonymization, data masking, and generalization help protect user identities while still allowing the data to be used for analysis. This approach is particularly useful when sharing datasets for research, analytics, or machine learning tasks without exposing sensitive personal details.
4. **Audit and Monitoring** – Tracks data access and modifications to detect suspicious activities and maintain compliance with organizational policies and regulations. Logging mechanisms record detailed information about user actions, system activities, and data usage. Continuous monitoring tools analyze these logs to detect anomalies, unauthorized access attempts, or potential security breaches, enabling quick response to threats.
5. **Network Security** – Protects the communication channels used to transfer data between distributed systems. Secure protocols such as encrypted communication channels help prevent attackers from intercepting or tampering with data during transmission. Firewalls and intrusion detection systems are also used to monitor network traffic and block malicious activities.
6. **Data Backup and Recovery** – Ensures that data can be restored in case of accidental deletion, hardware failure, or cyberattacks such as ransomware. Regular backups of critical datasets help maintain system availability and minimize data loss in case of unexpected failures.
7. **Compliance and Governance** – Organizations must follow data protection regulations and establish policies for proper data management. Data governance frameworks ensure that data is handled responsibly, securely, and ethically throughout its lifecycle. By applying these techniques, organizations can safely store, process, and share large-scale datasets while maintaining confidentiality, integrity, and compliance with data protection laws. Implementing strong security and privacy mechanisms ensures that big data analytics can



be performed effectively without compromising sensitive information. As the volume of data continues to grow, the importance of robust big data security frameworks will continue to increase in order to protect individuals, organizations, and digital infrastructures. Tools analyze these logs to detect anomalies, unauthorized access attempts, or potential security breaches, enabling quick response to threats.

8. Network Security – Protects the communication channels used to transfer data between distributed systems. Secure protocols such as encrypted communication channels help prevent attackers from intercepting or tampering with data during transmission. Firewalls and intrusion detection systems are also used to monitor network traffic and block malicious activities.
9. Data Backup and Recovery – Ensures that data can be restored in case of accidental deletion, hardware failure, or cyberattacks such as ransomware. Regular backups of critical datasets help maintain system availability and minimize data loss in case of unexpected failures.
10. Compliance and Governance – Organizations must follow data protection regulations and establish policies for proper data management. Data governance frameworks ensure that data is handled responsibly, securely, and ethically throughout its lifecycle.

By applying these techniques, organizations can safely store, process, and share large-scale datasets while maintaining confidentiality, integrity, and compliance with data protection laws. Implementing strong security and privacy mechanisms ensures that big data analytics can be performed effectively without compromising sensitive information. As the volume of data continues to grow, the importance of robust big data security frameworks will continue to increase in order to protect individuals, organizations, and digital infrastructures. The technical implementation of security measures is inseparable from the complex web of legal and ethical obligations. This experiment provided practical insight into how regulations like GDPR, CCPA, and HIPAA translate into concrete technical requirements. For instance, the "right to be forgotten" under GDPR necessitates not just the ability to delete a user's record from a primary database, but from all backups, data lakes, and derivative datasets a significant technical challenge in big data environments. Furthermore, the experiment touched upon the emerging issue of data sovereignty, where data is subject to the laws of the country in which it is physically stored. For organizations using distributed cloud infrastructure, this requires a clear understanding of where data resides at all times and the implementation of geo-fencing controls to prevent data from crossing borders in



violation of regulations. Beyond legal compliance, the ethical dimension of data use was considered, prompting a discussion on how organizations can move beyond merely checking boxes to genuinely respecting user privacy as a core principle. While technological controls form the backbone of data security, their effectiveness is ultimately determined by the people who use them. This experiment underscored that human error remains one of the largest vulnerabilities in any system. Even the most sophisticated encryption and access controls can be undermined by a single successful phishing attack or an employee's unintentional mishandling of data. Therefore, a comprehensive security strategy must prioritize ongoing security awareness training. This includes educating employees on how to recognize social engineering attempts, the importance of strong password hygiene, and the correct procedures for sharing and accessing sensitive information. By fostering a "security-first" culture, organizations transform their workforce from a potential liability into a proactive first line of defense, capable of identifying and reporting suspicious activities before they escalate into full-blown incidents.

Algorithm / Steps:

1. Import required libraries and distributed frameworks (e.g., Apache Spark).
2. Load large dataset from distributed storage such as Apache Hadoop (HDFS) or cloud storage.
3. **Assess Security Risks** – Evaluate potential threats such as unauthorized access, data breaches, or insider misuse.
4. **Apply Transformations** – Use RDD/DataFrame transformations such as map(), filter(), select(), groupBy() to manipulate data.
5. **Implement Privacy-Preserving Measures** – Apply techniques like data anonymization, masking, or encryption to protect sensitive information.
6. **Set Access Controls** – Define user roles, permissions, and authentication to restrict unauthorized access.
7. **Audit and Monitor** – Track data access, modifications, and sharing to ensure compliance and detect anomalies.
8. **Validate and Interpret** – Test security measures to ensure data remains protected while still usable for analysis.

Source Code:

```
# Import Libraries
import pandas as pd
from cryptography.fernet import Fernet
```



```
# Load Dataset
```

```
data = pd.read_csv("sensitive_data.csv")
```

```
# Identify Sensitive Column
```

```
sensitive_col = data['SSN']
```

```
# Generate Encryption Key
```

```
key = Fernet.generate_key()
```

```
cipher = Fernet(key)
```

```
# Encrypt Sensitive Data
```

```
data['SSN_encrypted'] = sensitive_col.apply(lambda x: cipher.encrypt(str(x).encode()).decode())
```

```
# Masking Example (show only last 4 digits)
```

```
data['SSN_masked'] = sensitive_col.apply(lambda x: "XXX-XX-" + str(x)[-4:])
```

```
# Display Secured Data
```

```
print(data.head())
```

Conclusion:

This experiment demonstrates how to identify and mitigate data security and privacy risks in big data environments. By implementing techniques such as encryption, masking, and anonymization, sensitive information can be protected while still allowing data analysis. Access controls, auditing, and monitoring ensure that only authorized users can access the data, maintaining confidentiality and integrity. These measures are crucial for complying with data protection regulations and building trust when handling large-scale sensitive datasets.

The practical implementation highlighted the layered approach required for comprehensive data security. Techniques like data masking help protect personally identifiable information (PII) during testing and development, while encryption safeguards data both at rest and in transit. Role-based access controls and detailed audit logs provide accountability and help detect unauthorized access attempts. As data breaches become more common and regulations like GDPR and HIPAA impose strict requirements, understanding these security measures is essential for organizations handling sensitive data. This experiment reinforces that security must be integrated throughout the data lifecycle, not added as an afterthought.

Furthermore, we explored the challenges unique to big data security, including the velocity, variety, and volume of data that make traditional security approaches insufficient. Distributed storage across multiple nodes increases the attack surface, requiring encryption at the file and network levels. Anonymization techniques such as k-anonymity and differential privacy were examined as methods to derive insights from data without exposing individual records. The experiment also demonstrated how security policies must balance protection with accessibility, ensuring that security measures do not hinder legitimate data analysis and business operations. By implementing these comprehensive security strategies, organizations can confidently leverage big data while maintaining customer trust and regulatory compliance in an increasingly data-driven world.

The experiment also highlighted the importance of data governance frameworks in establishing clear policies for data classification, handling, and retention. We learned that security is not solely a technical challenge but also involves organizational processes and employee training. Regular security audits and vulnerability assessments help identify potential weaknesses before they can be exploited. Additionally, incident response planning ensures that organizations can quickly



detect, contain, and recover from security breaches, minimizing damage and restoring normal operations. As cyber threats continue to evolve in sophistication, staying current with security best practices and emerging technologies like homomorphic encryption and secure multi-party computation becomes increasingly important for protecting sensitive data in big data environments.

Questions

1. What are the main security and privacy risks in big data environments?
2. How does data encryption help in protecting sensitive information?
3. What is the difference between data masking and data anonymization?
4. Why are access controls and auditing important in big data systems?
5. How can privacy-preserving techniques be applied without affecting data analysis?

PCU